

2

A data generation & validation pipeline

Months 2-3

WHAT YOU WILL DO

Design and build a reproducible pipeline for producing training data — including synthetic instructions, dialogues, and reasoning traces — and for validating its quality: deduplication, quality filtering, factual checks, model-as-judge review, and human sampling. Dataset versioning and traceability from day one.

WHAT YOU WILL HAVE BUILT

An operating data pipeline with quality metrics and versioned datasets, ready to feed training.

MATERIALS

BOOK — in progress

Ch. 10 Data is the cheapest efficiency lever

Ch. 25 Synthetic data & distillation

FROM PHASE 1 YOU REUSE

The BPE stepper's compression lens: better data → fewer, denser tokens → compute saved