

4

A benchmarking pipeline

Months 4–5

WHAT YOU WILL DO

Build the benchmarks that measure what your users actually need — knowledge and reasoning in the domains that matter — and automate the evaluation: a benchmark harness, per-capability metrics, judge-based evaluation, and reports comparable across models and versions.

WHAT YOU WILL HAVE BUILT

A benchmark suite per domain and an automated evaluation pipeline any model can be run through.

MATERIALS

BOOK — planned

Ch. 9 How you'll know it's working
(benchmarks, contamination, cheap proxies,
building an eval suite you trust)

GROUNDWORK ALREADY LIVE

Ch. 2 §2.4 — loss, perplexity, and reading curves:
the health metrics everything else builds on