

5

A living ranking of open models

Months 5–6

WHAT YOU WILL DO

Close the loop: a pipeline that periodically evaluates the leading open models — Llama, Qwen, Mistral, DeepSeek, Gemma, and whatever ships next — against your benchmark suite, weighing capability, inference cost, latency, license, and ease of deployment.

WHAT YOU WILL HAVE BUILT

A living model ranking with a dashboard and actionable model-selection guidance, maintained by the team that built it.

MATERIALS

BOOK — planned

Ch. 7 Scaling laws for training & inference

GROUNDWORK ALREADY LIVE

Ch. 13 Table 13.1 — attention at the frontier, model by model

The parameter counter — price any candidate model from its config in seconds